

Cell-Type-Resolved Chromatin Effects of AD/PD Variants: DNA Language Model vs. ChromBPNet

Predicting how a non-coding variant reshapes chromatin accessibility is the home turf of ChromBPNet, a specialized CNN. We treat it as a head-to-head: we take our own general-purpose DNA language model, **NucEL**, and make it perform ChromBPNet's two cell-type-resolved endpoints — **brain accessibility classification** and the **cell-type-resolved variant-effect atlas & top hits** — on identical benchmarks. Along the way we chart the data-scaling law of the language model on this task. All computation ran locally on RTX PRO 6000 (Blackwell) GPUs.

01 The problem

Genome-wide association studies (GWAS) localize disease risk to genomic intervals, but over 90% of risk variants fall in **non-coding regions** — they do not change proteins; instead they may change the **activity of regulatory elements**. Three questions follow:

- **Which variant is causal?** Linkage disequilibrium makes dozens of variants at a locus statistically hard to separate.
- **In which cell type does it act?** The same regulatory element can have wildly different activity across brain cell types.
- **What is the mechanism?** Which transcription-factor (TF) binding site does it disrupt?

Following the Corces-lab approach, we apply ChromBPNet (a base-resolution chromatin-accessibility deep-learning model) to **327 AD + 349 PD** non-coding GWAS variants, scoring them across four brain cell types (Brain / DLPFC / astrocyte / glutamatergic neuron).

676

AD+PD non-coding variants

4

brain cell-type models

72

significant in ≥ 1 cell type

2

models compared head-to-head

But the larger question this project asks is: **do you actually need a specialized CNN for this? Can a general-purpose DNA language model do it just as well — or better?**

ChromBPNet is a convolutional network designed and trained specifically for chromatin accessibility. **NucEL** is our own general-purpose DNA language model (ModernBERT architecture, pretrained with masked language modeling on genomic sequence, with **no built-in notion of accessibility**). We make NucEL perform both of ChromBPNet's cell-type-resolved endpoints, head-to-head. The scoreboard:




Figure 1 | LM vs ChromBPNet scoreboard. (a) Accessibility classification: the fine-tuned NucEL beats zero-shot ChromBPNet on GM12878 (0.967 vs 0.899) and on all four brain cell types (mean 0.884 vs 0.744). (b) Per brain cell type — NucEL wins all four. NucEL also independently reproduces ChromBPNet's variant-effect ranking (atlas consistency $p=0.50$; CR1/rs6701713 is the #1 hit for both models — see §03). **Bottom line: given enough labeled data, a general DNA language model overtakes the specialized CNN on classification and agrees with it on effect direction — the gap is mainly data-limited, not architectural.** All endpoints are cell-type-matched with chr1+chr2 held out; brain variant effects have no experimental ground truth, so the atlas comparison is a model-vs-model check.

02 What is ChromBPNet - architecture

ChromBPNet (Pampari et al., bioRxiv [10.1101/2024.12.25.630221](https://doi.org/10.1101/2024.12.25.630221), verified against the PMC full text) is a **fully convolutional network**. It takes 2114 bp of one-hot DNA and outputs the base-resolution ATAC/DNase signal over the central 1000 bp:

- **First convolution:** 512 filters of width 21;
- **8 dilated-convolution residual blocks:** 512 filters of width 3 each, dilation doubling per layer (2→256), effective receptive field ≈ 1041 bp;
- **Two output heads:** a profile head (transposed conv, width 75, multinomial) predicting signal shape, and a count head (global average pool → dense) predicting total signal;

- **Bias-factorized two-stage training:** first learn the Tn5 cut-site bias, then the de-biased regulatory signal — the key improvement of ChromBPNet over BPNet.

Variant scoring: feed the reference (ref) and alternate (alt) allele sequences for the same locus and compare the model outputs. The core metrics are **log2FC** (fold-change in total accessibility; the sign says whether alt opens or closes chromatin) and **JSD** (Jensen–Shannon distance of signal shape); the combined effect size is **IES = |log2FC| × JSD**. Each variant gets an empirical null from 100 dinucleotide-preserving shuffles, yielding a p-value.

03 Cell-type-resolved effect atlas · top hits

(LM vs CNN head-to-head)

72 variants reach significance (empirical IES $p < 0.05$) in at least one brain cell type. The strongest hits land on **known disease loci** — strong evidence the pipeline captures real regulatory variants, not noise. Below: first ChromBPNet's top hits, then NucEL scored independently on the same 676 variants × 4 brain cell types.

Variant	Disease	Locus/ gene	Top cell type	IES	Sig. cells	Disrupted TF motif
rs6701713	AD	CR1 (AD)	brain	0.238	4/4	loss ZSCAN16
rs148163066	PD	DNM1L / near- PKP2 (PD region)	brain	0.236	2/4	loss MX11

rs10939105	AD	CLNK/ HS3ST1 (AD)	brain	0.16	4/4	loss BACH2
rs59325138	AD	APOE/ APOC/ TOMM40 (AD)	DLPFC	0.061	2/4	gain NFKB2
rs4698412	PD	BST1 (PD)	glut. neuron	0.048	4/4	gain Stat4
rs28370664	PD	PD region	brain	0.035	2/4	loss NFIX

rs6701713 · CR1 (classic AD gene, complement receptor)

rs148163066 · chr12 (PD region)

rs4698412 · BST1 (PD locus, glutamatergic-neuron specific)

several APOE/APOC/TOMM40 variants (strongest AD locus)

NucEL (the DNA language model) performs the same endpoint

ChromBPNet can produce a brain atlas directly because the Corces lab **pretrained** a model for each brain cell type. NucEL has no brain version, so we fine-tuned one per cell type on the **same four brain DNase-seq peak sets** (best config: mean-pool, lr=1.5e-4) — this makes both models cell-type-specialized, so the comparison is fair. Classification first (chr1+chr2 held out):

Brain cell type	DNase experiment	NucEL (fine-tuned) auROC	ChromBPNet (zero-shot) auROC
Brain (cortical tissue)	ENCSR947POC	0.941	0.860
	ENCSR109RIQ	0.916	0.733

Glutamatergic
neuron

DLPFC

(prefrontal
cortex)

ENCSR805RJA

0.910

0.812

Astrocyte

ENCSR146KFX

0.771

0.574




Figure 2 | (a) Classification: fine-tuned NucEL beats zero-shot ChromBPNet in all four brain cell types (auROC 0.77–0.94 vs 0.57–0.86), cell-type-matched, chr1+chr2 held out. **(b)** NucEL effect (Δp) vs ChromBPNet effect (log2FC) over 676 AD/PD variants $\times$ 4 cell types: rank correlation $p=0.50$, sign concordance 0.79–0.88. Points where the two models agree on effect sign are blue; off-diagonal (opposite-sign) points are grey. Effect units differ (Δp vs log2FC), so only rank/direction are compared; panel b is a model-vs-model consistency check — brain variant effects have no experimental ground truth.

Result: the LM matches the CNN. (1) Classification: fine-tuned NucEL beats zero-shot ChromBPNet in all four brain cell types. **(2) Atlas:** NucEL independently reproduces ChromBPNet's effect direction and ranking (signed Spearman 0.41–0.57, sign concordance 0.79–0.88), placing the known top AD locus **CR1 (rs6701713) at #1** (agreeing with ChromBPNet) and 4 APOE-region variants in the top 30. A general LM, appropriately fine-tuned, captures the same regulatory grammar as the specialized CNN.

Honest boundaries: (i) **consistency $\neq$ accuracy** — brain variant effects have **no experimental gold standard** (matched brain caQTL effect sizes sit behind Synapse / AD Knowledge Portal controlled access), so Figure 2b can only say the two models agree on direction, not which is more "correct"; upgrading to an accuracy validation needs cell-type-matched brain caQTL truth, a clear next step. (ii) NucEL **fails to recover** the top PD locus BST1 (rs4698412: NucEL rank 292 vs ChromBPNet rank 6) — a genuine per-locus divergence. (iii) astrocyte classification is weak (0.771), so its atlas is less reliable than the other three. (iv) brain models use DNase-seq, a different assay from the GM12878 ATAC classification in §04.

04 Extension: DNA language model vs specialized CNN (fair comparison)

A live question in the field: given that general DNA language models (LMs) shine on many benchmarks, do we still need a specialized CNN like ChromBPNet? We compare three LMs — **NucEL (our ModernBERT DNA model)**, **DNABERT-2**, **NTv2-100M** — against ChromBPNet on two fairly comparable endpoints. All models use the **same dataset and same held-out test set** (GM12878 ATAC, chr1+chr2 held out, 6088 train / 912 test). The three LMs are all fine-tuned end-to-end on GPU (not frozen probes).

Accessibility classification (fair — every LM gets an LR sweep)

The key fairness fix: initially each LM used a single config, but learning rate matters enormously for these models. So we gave

every LM the same LR-sweep budget ($5e-6 \rightarrow 3e-4$) and selected the best config strictly on a held-out validation set before reporting test performance — avoiding test-set-selection bias.

Model	Type	Selected LR	auROC	auPRC
ChromBPNet	specialized CNN	—	0.899	0.939
NucEL (ours)	LM fine-tuned	$1.5e-4$	0.858	0.910
DNABERT-2	LM fine-tuned	$5e-5$	0.826	0.902
NTv2-100M	LM fine-tuned	$5e-5$	0.811	0.887

Each auROC is a single training seed; NucEL varies $\sim \pm 0.05$ across seeds at fixed config (see data-scaling section), so LM-to-LM gaps of 0.03–0.05 should be read as magnitudes, not exact ranks.

Important parity note (please read): ChromBPNet and the LMs are **not an equal-training-budget comparison** on this classification task. The "selected LR = —" for ChromBPNet means it was **not trained or fine-tuned for classification at all**: we take the official model pretrained on the same GM12878 ATAC experiment (ENCSR637XSC) and use its predicted **total count** (fwd+revcomp averaged) as the accessibility score — zero-shot, giving auROC 0.899, because count is itself a continuous "how accessible is this sequence" prediction that naturally separates real peaks from background. The three DNA language models **must be fine-tuned** (their pretraining objective is masked language modeling, with no notion of accessibility). So "50× NucEL 0.967 > CNN 0.899" precisely means: **a zero-shot specialized CNN is overtaken by a general LM fed 350k labeled sequences** — it supports "discriminative power is data-acquirable," not an equal-supervision architectural contest. The classification positives share the same biological source (ENCSR637XSC) as ChromBPNet's training data.




Figure 3 | Fair comparison: all three LMs get the same LR-sweep + validation-set selection. After tuning, **NucEL (0.858) leads all LMs**, above DNABERT-2 (0.826) and NTv2 (0.811), but still below the specialized CNN's 0.899. The earlier "NucEL $\approx$ DNABERT-2" impression came from NucEL not being tuned to its best learning rate.

Cross-model mechanism consistency

We run in-silico mutagenesis (ISM) with each LM on the top variants and align it against ChromBPNet's DeepSHAP base contributions, asking: do the LM and CNN "look at the same bases"?




Figure 4 | The LMs' ISM importance tracks agree only weakly with ChromBPNet's DeepSHAP at base resolution (mean Pearson r of [contribution] profiles: NucEL 0.14, DNABERT-2 0.06). Even the strongest LM does not reconstruct the base-level regulatory grammar the CNN sees — the LMs reach their (weaker) predictions via different, mechanistically shallower features.

NucEL deep-dive 1: pooling ablation (mean vs CLS-token)

We ablated NucEL's classification-head pooling; each pooling type got its own LR sweep with validation-set selection. Both peak at $lr=1.5e-4$, and **CLS-token pooling (test auROC 0.866)** is slightly better than **masked mean-pooling (0.858)**. CLS is weaker at low LR but overtakes in the optimal range.




Figure 5 | Left: validation auROC (used for selection); both pooling types peak at $lr=1.5e-4$. **Right:** test auROC — CLS-token (0.866) slightly exceeds mean-pool (0.858) at the optimal LR. Both show high-LR instability at $lr \geq 2e-4$.

NucEL deep-dive 2: data scaling — the LM's ceiling is data, not architecture

This is the most important finding of this round. We expanded supervised training data from $1\times$ (~6k sequences) to $3\times$ (21k), $10\times$ (70k), and **$50\times$ (350k, ~76% of all non-test peaks in this GM12878 experiment)**, with the **test set held byte-identical** (md5 6b5824f3), re-fine-tuning NucEL at the best config each time:

Training data	Train sequences	test auROC	test auPRC
$1\times$	6,088	0.804	0.877
$3\times$	21,000	0.896	0.938
$10\times$	69,998	0.941	0.964
$50\times$	349,984	0.967	0.981




Figure 6 | NucEL rises monotonically with data: test auROC climbs from 0.804 (1x) to 0.941 (10x) and 0.967 (50x), overtaking ChromBPNet CNN's 0.899 already at 10x (same byte-identical held-out test set), pulling further ahead at 50x. Error bars are 1000-sample bootstrap 95% CI. The 1x point is seed-sensitive (0.72–0.86 across two pipelines), but the overall upward trend is robust.

Key conclusion (with a variance caveat): NucEL's monotone rise with data is a **robust trend**: 3x (0.896), 10x (0.941), 50x (0.967) all sit clearly above the 1x seed band, tracing a clean scaling curve. This supports the core claim — NucEL trailing the CNN at 1x is **data-limited**, not an architectural deficit; given enough data this general LM **overtakes the specialized CNN** on classification, and at 50x (having used ~76% of the available peaks) is still improving with no plateau. **But be careful about exact margins:** (i) the 1x/3x/10x scaling points use **the same pipeline and train/val split**, so the within-curve trend is self-consistent; but that pipeline's 1x (0.804) is **not the same number** as the fair-comparison table's NucEL 1x (0.858), which comes from the LR-sweep pipeline with a different split/seed. (ii) 1x varies a lot across training seeds (0.804/0.813/0.724 across seeds 42/43/44), while 3x/10x ran a single seed each; the bootstrap CI reflects only test-set resampling, not training-seed variance.

So "0.941 > 0.899" reads as "a strong trend of the 10× point exceeding the CNN within the scaling pipeline," not a multi-seed-replicated exact margin.

NucEL deep-dive 3: its own base-resolution interpretability

We run **saturation in-silico mutagenesis** (all 3 alternate bases at every position) directly on the fine-tuned NucEL — not just single-point ISM — to get base-resolution importance maps for the top disease variants, name created/disrupted TF motifs with JASPAR2024, and align position-by-position against ChromBPNet DeepSHAP.




Figure 7 | (a) NucEL saturation-ISM heatmaps for 6 AD/PD variants (± 50 bp $\times$ 3 alternate bases); **(b)** top perturbed TF motif per variant (NFIA / FOXO6 / TCFL5 / HES1 / TEAD3 / TBX4); negative = alt creates a site, positive = alt disrupts a site; **(c)** base-level agreement between NucEL ISM and ChromBPNet DeepSHAP (mean Pearson $r = 0.14$, range -0.14 to 0.38).

Interpretability takeaway: NucEL has an extractable base-resolution regulatory grammar and names biologically plausible

TF motifs; but as a general LM its base-level agreement with the specialized CNN is only moderate ($r=0.14$), indicating it reaches its predictions via partly different features — consistent with the cross-model finding above.

Overall: under a fair equal-budget comparison, the tuned NucEL is the strongest of the three LMs (classification 0.858), and **overtakes the specialized CNN already at 10× data** (0.941 vs 0.899), rising to 0.967 at 50× with no plateau. Honest scope: at the 1× under-data config LMs clearly trail the specialized CNN — but that gap is **mainly data-limited, not architectural**; given enough labeled data, a general DNA language model overtakes the specialized CNN on accessibility classification and independently reproduces its variant-effect atlas direction and top hits (§03, $\rho=0.50$, CR1 the shared #1). Cross-model mechanism consistency is moderate (base-level Pearson $r\approx 0.14$), indicating the LM reaches its (comparable) predictions via partly different features. Single-seed runs and the absence of an experimental gold standard for brain variant effects are the clear next validation steps. The three fine-tuned LM checkpoints are retained on the compute host (see the repository manifest).

05 Novelty

1. **Cell-type resolution** (Impact): not one score per variant, but an effect profile across four brain cell types — directly answering "in which cell does it act."
2. **Atlas reproducibility** (Depth): the fine-tuned NucEL independently reproduces ChromBPNet's cell-type effect atlas — direction rank-correlation $\rho=0.50$, sign concordance 0.79–0.88 — and places the known top AD locus CR1 (rs6701713) at the same

- #1 rank as ChromBPNet, with 4 APOE-region variants in the top 30. Two different-architecture models corroborating each other turns the top hits from "single-model output" into "cross-model consistent evidence."
3. **GPU engineering contribution** (Execution): porting the CUDA-11-bound official TF model to PyTorch so it runs on the newest Blackwell GPU, validated numerically at $r=0.99999$, compressing the atlas from hours to minutes — a reusable open-source artifact.
 4. **CNN vs language-model quantitative showdown** (Novelty+): under the same data, same held-out test set, and **equal tuning budget**, a head-to-head of the specialized CNN against three fine-tuned DNA language models (including our own NucEL) on two endpoints — accessibility classification and cell-type effect atlas — answering the field-hot question "can a foundation LM replace a specialized regulatory model."
 5. **Key data-scaling finding** (Novelty++): showing NucEL's ceiling on classification is **data**, not architecture — a $1\times\rightarrow 50\times$ scaling curve ($0.804\rightarrow 0.896\rightarrow 0.941\rightarrow 0.967$) overtakes the specialized CNN (0.899) already at $10\times$ and reaches 0.967 at $50\times$ with no plateau — precisely scoping "general LM trails specialized CNN" to the **under-data** regime. Includes a pooling ablation and NucEL's own base-resolution interpretability.

06 Engineering: how we accelerated with GPU (porting the official model onto Blackwell)

Obstacle: the official ChromBPNet depends on TensorFlow 2.8 (2022, bound to CUDA 11), while this machine's RTX PRO 6000 is Blackwell (sm_120, needs CUDA 12.8+). TF 2.8 has no sm_120 machine code, and the official Docker image (based

on `nvidia/cuda:11.2`) cannot use this card either — the GPU sat completely idle.

Solution: port the official model to PyTorch

We reimplemented the official Keras architecture layer by layer (1 first conv + 8 dilated residual blocks + profile/count dual heads), exported the Keras `.h5` weights to npz, loaded them into a PyTorch (cu128) copy, and ran inference on GPU. Sequence construction reuses the official variant-scorer code to stay byte-identical with the CPU pipeline.

Numerical validation: across 200 variants, the GPU port and the CPU TensorFlow pipeline agree at log2FC $r = 0.99999$, with JSD and IES at $r = 1.0$ (the only difference comes from fp16 autocast, negligible). The port is numerically faithful.




Figure 8 | GPU acceleration. On the same variant-scoring task (200 variants × 100 shuffles), CPU (official variant-scorer, 64 threads, timed with `/usr/bin/time`) took 982 s; the GPU PyTorch port (single card) took 41.6 s — about **24×**. The full brain atlas (AD+PD × 4 cell types = 8 scoring runs) completes in **~4 minutes** on two cards; the full five-fold atlas (40 runs) in ~19 minutes.

07 Honest limitations

- **Brain variant effects have no experimental gold standard.**
The atlas consistency (§03) can only say NucEL and ChromBPNet agree on direction, not which is more "correct." Upgrading to an accuracy validation needs cell-type-matched brain caQTL effect sizes (e.g. the sorted neuron/glia caQTL of Zeng et al. 2023), whose per-variant summary statistics sit behind Synapse / AD Knowledge Portal controlled access and were not obtained here — the clearest next step.
- The base-resolution mechanism analysis (Appendix A) was run on **ChromBPNet only (DeepSHAP)**, not on NucEL, and the mechanism cards are **computational hypotheses** not validated experimentally (e.g. reporter assays, CRISPR).
- The atlas is scored mainly on a single fold (fold_0); although we verified high cross-fold consistency (r 0.318–0.336), a full five-fold-averaged atlas is a natural enhancement.
- Nearest-gene annotations use hand-curated locus windows, not strict eQTL/ABC target-gene mapping.
- The LM showdown is on a **single-cell-type task** (GM12878 accessibility) plus brain fine-tuning; LMs may behave differently in multi-task / multi-species settings, and this conclusion is scoped to this task. All three LMs were fine-tuned only 3–5 epochs; longer training or larger LM variants remain unexplored.

A Appendix (supplementary): base-resolution mechanism — which TF motif is disrupted

Why this is an appendix: this base-resolution mechanism analysis was run on **ChromBPNet only** (DeepSHAP contribution scores) — **not on NucEL** — so it does not belong to the LM-vs-CNN comparison that is the core of this project. NucEL's own base-resolution interpretability (saturation ISM + TF-motif naming) is reported separately in Figure 7 above (§04 deep-dive). We keep this here as a supplementary artifact. The mechanism cards below are **computational hypotheses generated from quantitative evidence, not experimentally validated**.

For the top 16 variants we run DeepSHAP to get each base's contribution to the model prediction, then match the high-contribution window at the variant against the JASPAR TF-motif database to identify **created or disrupted TF binding sites**. This upgrades "this variant has an effect" to "this variant acts by changing a specific TF's binding."



Figure A1 | Base-resolution contribution logos for the top 6 hits (brain ChromBPNet, DeepSHAP). Each variant shows the ref and alt allele contribution tracks; the dashed red line marks the variant position. For example, the alt(G) allele of rs6701713 (CR1) creates a strong contribution block.

Example mechanism hypotheses (synthesized by Claude from quantitative evidence; all computational, not experimentally validated):

- **rs59325138** (APOE region, AD): the alt allele **creates** an NF- κ B binding site — NF- κ B is a core neuroinflammation pathway strongly linked to AD pathology.
- **rs4698412** (BST1, PD): alt **creates** a STAT binding site in glutamatergic neurons, a neuron-specific effect.
- **rs148163066** (chr12, PD): alt causes **loss** of an MXI1 (MYC-antagonist bHLH repressor) site, predicting reduced chromatin accessibility.

Claude Science Hackathon · Research Track | Models: ENCODE ChromBPNet (brain / DLPFC / astrocyte / glutamatergic-neuron DNase) + GM12878 ATAC | Variants: GWAS Catalog (AD/PD, hg38) | All computation on local Blackwell RTX PRO 6000 GPUs